

rraftec: A Deep Reinforcement Learning agent able to play the game of Crafter

Rares FOLEA

Computer Science & Engineering Department
Faculty Of Automatic Control And Computers
University Politehnica Of Bucharest
Bucharest, Romania
rares.folea@stud.acs.upb.ro

Andrei VATAVU

Computer Science & Engineering Department
Faculty Of Automatic Control And Computers
University Politehnica Of Bucharest
Bucharest, Romania
andrei.vatavu@stud.acs.upb.ro

Abstract—rraftec: A Deep Reinforcement Learning agent able for the game of Crafter [1].

Index Terms—Deep Reinforcement Learning, Crafter, Agent, DQN, Double-DQN, pytorch

CONTENTS

I	Introduction	1
II	Implementation of rraftec	1
II-A	Neural NetV1	1
II-B	Neural NetV2	2
III	Training of rraftec	2
IV	Evaluating results of rraftec against Crafter	2
V	Appendix A: 4-seed DDQN rraftec-agents trained 1KK steps	3
VI	Appendix B: 4-seed DQN rraftec-agents trained 1KK steps	3

I. INTRODUCTION

Crafter features randomly generated 2D worlds where the player needs to forage for food and water, find shelter to sleep, defend against monsters, collect materials, and build tools. **Crafter** aims to be a fruitful benchmark for reinforcement learning [1].

In our implementation, we did a benchmark between the performance of a random agent and two Deep Reinforcement Learning agents, based using **DQN**(Deep Q-learning) and **Double-DQN**.

Because Q-learning evaluates the future maximum approximated action value using the same Q function as the current action selection policy, it can occasionally overestimate the action values in noisy situations, slowing learning. To address this, a variation known as Double Q-learning has been developed. Double Q-learning[2] is




Figure 1: The possible maps of Crafter.




Figure 2: The Deep Convolutional NetworkV1 chose by us for the Deep Reinforcement Learning agent rraftec.

an off-policy reinforcement learning system in which the value evaluation policy differs from the policy used to pick the next action.

II. IMPLEMENTATION OF RRAFTEC

Our implementation for an agent able to play this game is available open-source, at:

<https://github.com/raresraf/rraftec>

We have tried two approaches for the neural network.

A. Neural NetV1

The architecture for the deep neural network that we have used is:

Listing 1: Implementation of the network using pyTorch

```
def get_estimator(action_num, device,
                 input_ch=4,
                 lin_size=32):
    return nn.Sequential(
        nn.Conv2d(input_ch, 16, kernel_size=3),
        nn.ReLU(inplace=True),
        nn.Conv2d(16, 4, kernel_size=3),
        nn.ReLU(inplace=True),
        View(),
        nn.Linear(4 * 80 * 80, 80 * 80),
        nn.ReLU(inplace=True),
        nn.Linear(80 * 80, lin_size),
        nn.ReLU(inplace=True),
        nn.Linear(lin_size, action_num),
    ).to(device)
```

B. Neural NetV2

We have also trained the model on a slightly different network, where we have seen a little improvement in terms of rewards.

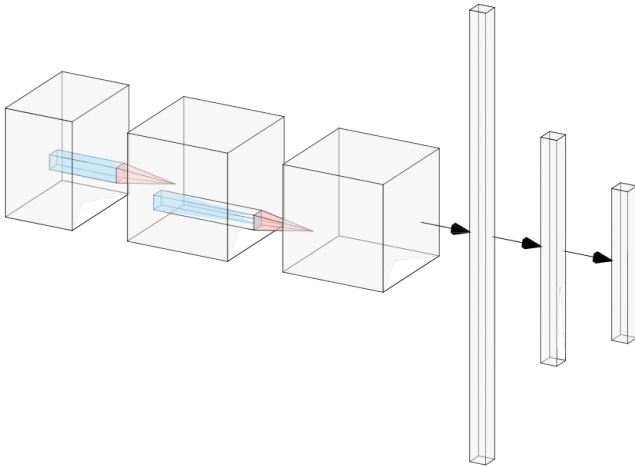


Figure 3: The Deep Convolutional NetworkV2 chose by us for the Deep Reinforcement Learning agent rraftec.

Listing 2: Implementation of the second network using pyTorch

```
def get_estimator(action_num, device,
                 input_ch=4,
                 lin_size=32):
    return nn.Sequential(
        nn.Conv2d(input_ch, 8, kernel_size=3),
        nn.ReLU(inplace=True),
        nn.Conv2d(8, 8, kernel_size=3),
        nn.ReLU(inplace=True),
        View(),
        nn.Linear(8 * 80 * 80, 8 * 80),
        nn.ReLU(inplace=True),
        nn.Linear(8 * 80, lin_size),
```

```
nn.ReLU(inplace=True),
nn.Linear(lin_size, action_num),
).to(device)
```

In our implementation, certain modules have been used from the previous labs of AAIT@UPB, containing implementations of the **DQN** and the **Double-DQN**.

Link to the resource:

<https://colab.research.google.com/drive/1B1sQXuyyTkfHza9Pw5kAUUaSEX2FWPmY>

III. TRAINING OF RRAFTEC

In the training, we have used the following *params* and *hyperparams* for **DQN** and **DDQN**:

A replay memory:

(size=1000, batch_size=32),

Adam optimizer for the network:

(learning rate=1e-3, eps=1e-4),

linear decay epsilon schedule:

(start=1.0, end=0.1, steps=100000),

Warmup steps: 10000,

Update steps: every 1 iteration.

We have trained the model using 4 seeds, each for **1000000** steps.

IV. EVALUATING RESULTS OF RRAFTEC AGAINST CRAFTER

We have evaluated every 10000 the reward obtained by the rraftec agent against Crafter. The total training took **1000000** steps, and at the end of the training, **Double-DQN** (clearly) have overall performed better than the initial *random* agent on the game of Crafter. **DQN** showed only some slightly improvements from the initial *random* agent on the game of Crafter, even though there were episodes in which **DQN** still performed worse than the random agent.

Double-DQN has clearly outclassed **DQN** in terms of reward obtained.

REFERENCES

- [1] Danijar Hafner. Benchmarking the spectrum of agent capabilities. *arXiv preprint arXiv:2109.06780*, 2021.
- [2] Hado Hasselt. Double q-learning. In J. Lafferty, C. Williams, J. Shawe-Taylor, R. Zemel, and A. Culotta, editors, *Advances in Neural Information Processing Systems*, volume 23. Curran Associates, Inc., 2010.




Figure 4: Training the Deep Reinforcement Learning agent rraftec using the DDQN strategy on the game of Crafter.

V. APPENDIX A: 4-SEED DDQN RRAFTEC-AGENTS TRAINED 1KK STEPS

This appendix shows 4 training episodes of the Deep Reinforcement Learning agent rraftec using the **Double-DQN** strategy and the NeuralNetV2.

VI. APPENDIX B: 4-SEED DQN RRAFTEC-AGENTS TRAINED 1KK STEPS

This appendix shows 4 training episodes of the Deep Reinforcement Learning agent rraftec using the **DQN** strategy and the NeuralNetV2.




Figure 5: Training the Deep Reinforcement Learning agent rraftec using the DQN strategy did not show good performances on the game of Crafter.


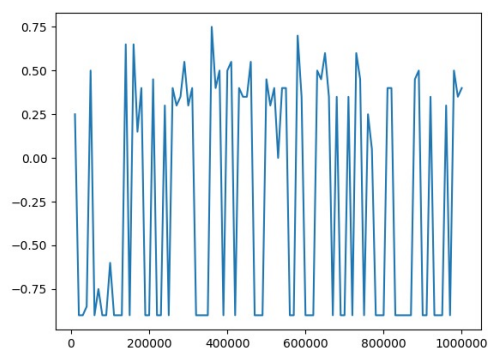
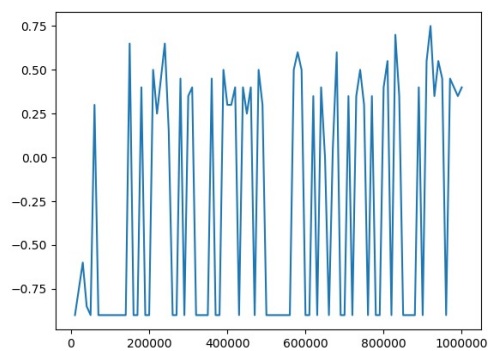

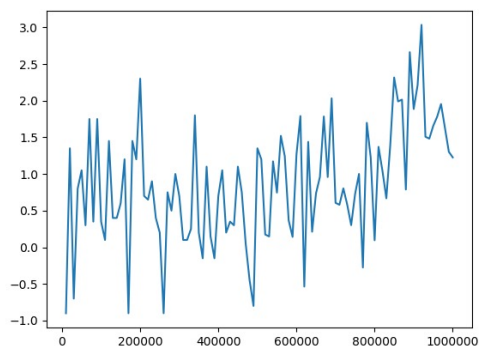
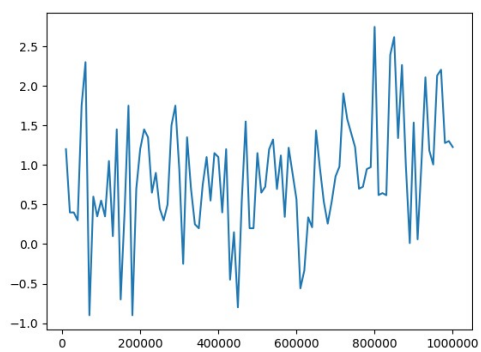
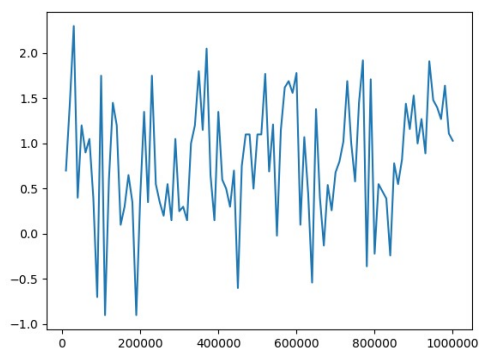
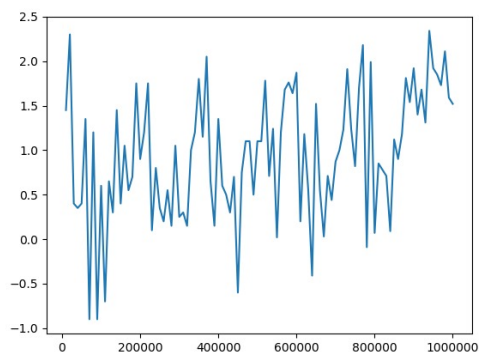


Figure 6: 4-seed training of the Deep Reinforcement Learning agent rraftec using the DDQN strategy.

Figure 7: 4-seed training of the Deep Reinforcement Learning agent rraftec using the DQN strategy.